

# ISL - Chapter 6 Lab Tutorials

## Linear Model Selection and Regularization

An introduction to Statistical Learning, with Applications in R  
- G. James, D. Witten, T. Hastie, R. Tibshirani

*Thu Nguyen*

*26 June, 2019*

### Contents

<b>6.5. Lab 1: Subset Selection Methods</b>	<b>2</b>
6.5.1. Best Subset Selection . . . . .	2
6.5.2. Forward and Backward Stepwise Selection . . . . .	5
6.5.3. Choosing among Models using Validation Set Approach and CV . . . . .	6
<b>6.6. Lab 2: Ridge Regression and the Lasso</b>	<b>8</b>
6.6.1. Ridge Regression . . . . .	8
6.6.2. The Lasso . . . . .	12
<b>6.7. Lab 3: PCR and PLS Regression</b>	<b>14</b>
6.7.1. Principal Components Regression . . . . .	14
6.7.2. Partial Least Squares . . . . .	15

---

### Main Contents:

1. Subset Selection
  2. Shrinkage Methods
  3. Dimension Reduction Methods
  4. Considerations in High Dimensions
-

## 6.5. Lab 1: Subset Selection Methods

---

### 6.5.1. Best Subset Selection

```
library(ISLR)
# Hitters dataset from ISLR, and remove NA from Salary column
Hitters <- na.omit(Hitters)
print(paste0('Dimension after removing NA: ', dim(Hitters)))
```

```
## [1] "Dimension after removing NA: 263" "Dimension after removing NA: 20"
```

`regsubsets()` Best subset selection, from library leaps

```
# regsubsets(): Best subset selection, from library 'leaps'
library(leaps)
regfit.full <- regsubsets(Salary ~ ., data = Hitters)
summary(regfit.full)
```

```
## Subset selection object
## Call: regsubsets.formula(Salary ~ ., data = Hitters)
## 19 Variables (and intercept)
##           Forced in Forced out
## AtBat      FALSE      FALSE
## Hits       FALSE      FALSE
## HmRun       FALSE      FALSE
## Runs       FALSE      FALSE
## RBI        FALSE      FALSE
## Walks      FALSE      FALSE
## Years      FALSE      FALSE
## CAtBat     FALSE      FALSE
## CHits      FALSE      FALSE
## CHmRun     FALSE      FALSE
## CRuns      FALSE      FALSE
## CRBI       FALSE      FALSE
## CWalks     FALSE      FALSE
## LeagueN    FALSE      FALSE
## DivisionW  FALSE      FALSE
## PutOuts    FALSE      FALSE
## Assists    FALSE      FALSE
## Errors     FALSE      FALSE
## NewLeagueN FALSE      FALSE
## 1 subsets of each size up to 8
## Selection Algorithm: exhaustive
##           AtBat Hits HmRun Runs RBI Walks Years CAtBat CHits CHmRun CRuns
## 1 ( 1 ) " " " " " " " " " " " " " " " " " " " " " "
## 2 ( 1 ) " " "*" " " " " " " " " " " " " " " " "
## 3 ( 1 ) " " "*" " " " " " " " " " " " " " " " "
## 4 ( 1 ) " " "*" " " " " " " " " " " " " " " " "
## 5 ( 1 ) "*" "*" " " " " " " " " " " " " " " " "
## 6 ( 1 ) "*" "*" " " " " " " "*" " " " " " " " "
## 7 ( 1 ) " " "*" " " " " " " "*" " " "*" "*" " " "
```

```
## 8 ( 1 ) "*" " " " " " " "*" " " " " " " "*" "*"
##          CRBI CWalks LeagueN DivisionW PutOuts Assists Errors NewLeagueN
## 1 ( 1 ) "*" " " " " " " " " " " " "
## 2 ( 1 ) "*" " " " " " " " " " " " "
## 3 ( 1 ) "*" " " " " " " "*" " " " " " "
## 4 ( 1 ) "*" " " " " "*" "*" " " " " " "
## 5 ( 1 ) "*" " " " " "*" "*" " " " " " "
## 6 ( 1 ) "*" " " " " "*" "*" " " " " " "
## 7 ( 1 ) " " " " " " "*" "*" " " " " " "
## 8 ( 1 ) " " "*" " " "*" "*" " " " " " "
```

By default: `regsubsets()` only include up to 8 variables, to increase, specify `nvmax = p`, with  $p$  variables

```
regfit.full <- regsubsets(Salary ~ ., data = Hitters, nvmax = 19)
reg.summary <- summary(regfit.full)
mes <- 'Available attributes of model'
print(cat(mes, '\n', names(reg.summary), '\n'))
```

```
## Available attributes of model
## which rsq rss adjr2 cp bic outmat obj
## NULL
```

```
#  $R^2$ 
round(reg.summary$rsq, 4)
```

```
## [1] 0.3215 0.4252 0.4514 0.4754 0.4908 0.5087 0.5141 0.5286 0.5346 0.5405
## [11] 0.5426 0.5436 0.5445 0.5452 0.5455 0.5458 0.5460 0.5461 0.5461
```


**Interpretation:** 1 var:  $R^2 = .32$ , 2 vars:  $R^2 = .425$ , ... 19 vars:  $R^2 = .546$

**Plot**

```
# Plot
par(mfrow=c(2,2)); par(mar=c(3,5,1,1))
# RSS
plot(reg.summary$rsq, xlab = 'Number of Variables', ylab = 'RSS', type = 'l')
# Adjusted  $R^2$ 
plot(reg.summary$adjr2, xlab = 'Number of Variables', ylab = 'Adjusted RSq', type = 'l')
# Max(Adjusted  $R^2$ )
points(11, reg.summary$adjr2[11], col = 'red', cex = 2, pch = 20)

#  $C_p$ 
plot(reg.summary$cp, xlab = 'Number of Variables', ylab = 'Cp', type = 'l')
points(10, reg.summary$cp[10], col = 'red', cex = 2, pch = 20)

# BIC
plot(reg.summary$bic, xlab = 'Number of Variables', ylab = 'BIC', type = 'l')
points(6, reg.summary$bic[6], col = 'red', cex = 2, pch = 20)
```



```
attr.names <- c('Max(Adjusted R^2)', 'Min(Cp)', 'Min(BIC)')
attr.values <- c(which.max(reg.summary$adjr2), which.min(reg.summary$cp), which.min(reg.summary$bic))
data.frame(Attributes = attr.names, Number_of_Variables = attr.values)
```

```
##           Attributes Number_of_Variables
## 1 Max(Adjusted R^2)                11
## 2           Min(Cp)                 10
## 3           Min(BIC)                 6
```

```
# Plot with labelled best variables
# par(mfrow=c(1,1))
# plot(regfit.full, scale = 'r2')
# plot(regfit.full, scale = 'adjr2')
# plot(regfit.full, scale = 'Cp')
# plot(regfit.full, scale = 'bic')
```

## Interpretation:

- The top row of each plot contains a black square for each variable selected according to the optimal model. For instance, we see that several models share a BIC close to  $-150$ . However, the model with the lowest BIC is the 6-variable model that contains only AtBat, Hits, Walks, CRBI, DivisionW, and PutOuts.

```
mes <- 'Coefficients of model with 6 var'
print(cat(mes, '\n', round(coef(regfit.full, 6), 4), '\n'))
```

```
## Coefficients of model with 6 var
## 91.5118 -1.8686 7.6044 3.6976 0.643 -122.9515 0.2643
## NULL
```

```
round(coef(regfit.full, 6), 4)
```

```
## (Intercept)      AtBat      Hits      Walks      CRBI      DivisionW
## 91.5118      -1.8686      7.6044      3.6976      0.6430     -122.9515
## PutOuts
## 0.2643
```

## 6.5.2. Forward and Backward Stepwise Selection

to specify method: `method = 'forward/backward'`

### Forward

```
regfit.fwd <- regsubsets(Salary ~ ., data = Hitters, nvmax = 19, method = 'forward')
# summary(regfit.fwd)
```

### Backward

```
regfit.bwd <- regsubsets(Salary ~ ., data = Hitters, nvmax = 19, method = 'backward')
# summary(regfit.bwd)
```

### Comment:

- For best models with 1-6 vars, both are the same
- 7 vars: Forward and Backward return different models

### Full model

```
round(coef(regfit.full, 7), 2)
```

## (Intercept)	Hits	Walks	CAtBat	CHits	CHmRun
## 79.45	1.28	3.23	-0.38	1.50	1.44
## DivisionW	PutOuts				
## -129.99	0.24				

### Forward Stepwise model

```
round(coef(regfit.fwd, 7), 2)
```

## (Intercept)	AtBat	Hits	Walks	CRBI	CWalks
## 109.79	-1.96	7.45	4.91	0.85	-0.31
## DivisionW	PutOuts				
## -127.12	0.25				

### Backward Stepwise model

```
round(coef(regfit.bwd, 7), 2)
```

## (Intercept)	AtBat	Hits	Walks	CRuns	CWalks
## 105.65	-1.98	6.76	6.06	1.13	-0.72
## DivisionW	PutOuts				
## -116.17	0.30				

---

### 6.5.3. Choosing among Models using Validation Set Approach and CV

```
# Split into train and test sets
set.seed(1)
train <- sample(c(TRUE, FALSE), nrow(Hitters), rep = TRUE)
test <- (!train)
# Best subset model
regfit.best <- regsubsets(Salary ~ ., data = Hitters[train,], nvmax = 19)
# Validation set test error
test.mat <- model.matrix(Salary ~ ., data = Hitters[test,])
# Test set MSE
val.errors <- c()
for (i in 1:19) {
  coefi <- coef(regfit.best, id = i)
  pred <- test.mat[, names(coefi)] %*% coefi
  val.errors[i] <- round(mean((Hitters$Salary[test] - pred)^2), 0)
}
```

MSE on test set for different number of variables, arranged by row.

```
matrix(val.errors, ncol = 5, byrow = T)
```

```
##      [,1] [,2] [,3] [,4] [,5]
## [1,] 220968 169157 178518 163426 168418
## [2,] 171271 162377 157909 154056 148162
## [3,] 151156 151742 152214 157359 158541
## [4,] 158743 159973 159860 160106 220968
```

Model with the best MSE

```
# Model with min(MSE)
mes <- 'The number of variables giving lowest MSE on test set: '
print(paste0(mes, which.min(val.errors)))
```

```
## [1] "The number of variables giving lowest MSE on test set: 10"
```

```
coef(regfit.best, which.min(val.errors))
```

```
## (Intercept)      AtBat      Hits      Walks      CAtBat      CHits
## -80.2751499 -1.4683816  7.1625314  3.6430345 -0.1855698  1.1053238
##      CHmRun      CWalks      LeagueN      DivisionW      PutOuts
##  1.3844863 -0.7483170  84.5576103 -53.0289658  0.2381662
```

predict.regsubsets() function, and 10-fold Cross-Validation

```
# predict fn, by hand
predict.regsubsets <- function(object, newdata, id, ...) {
  form <- as.formula(object$call[[2]])
  mat <- model.matrix(form, newdata)
  coefi <- coef(object, id = id)
  xvars <- names(coefi)
  mat[,xvars] %*% coefi
}
```

```

}

# 10-fold Cross-Validation
k <- 10
set.seed(1)
folds <- sample(1:k, nrow(Hitters), replace = TRUE)
cv.errors <- matrix(NA, k, 19, dimnames = list(NULL, paste(1:19)))

```

Average MSE on test set for each number of variables after 10-fold CV

```

# Steps: prediction -> test set errors
# Loop 1: j for each fold
for (j in 1:k) {
  best.fit <- regsubsets(Salary ~ ., data = Hitters[folds != j,], nvmax = 19)
  # Loop 2: i for each variable
  for (i in 1:19) {
    pred <- predict(best.fit, Hitters[folds == j,], id = i)
    cv.errors[j,i] <- mean( (Hitters$Salary[folds == j] - pred)^2 )
  }
}
mean.cv.errors <- round(apply(cv.errors, 2, mean),0)
mean.cv.errors

```

```


##      1      2      3      4      5      6      7      8      9     10
## 160093 140197 153117 151159 146841 138303 144346 130208 129460 125335
##      11     12     13     14     15     16     17     18     19
## 125154 128274 133461 133975 131826 131883 132751 133096 132805

```

```

par(mfrow = c(1,1)); par(mar=c(2,4,1,1))
plot(mean.cv.errors, type = 'b')

```



**Comment:** based on plot, best model has 11 variables

Coefficients for the model

```

reg.best <- regsubsets(Salary ~ ., data = Hitters, nvmax = 19)
round(coef(reg.best, 11),2)

```

```

## (Intercept)      AtBat      Hits      Walks      CAtBat      CRuns
##      135.75      -2.13      6.92      5.62      -0.14      1.46
##      CRBI      CWalks      LeagueN      DivisionW      PutOuts      Assists
##      0.79      -0.82      43.11      -111.15      0.29      0.27

```

## 6.6. Lab 2: Ridge Regression and the Lasso

`glmnet()` from package `glmnet`, for Ridge Regression: `alpha = 0`, Lasso: `alpha = 1`, also, by default, `glmnet()` standardizes the variables automatically.

```
library(glmnet)
# Rmb to remove na/missing values
# Convert from data.frame to matrix, also, categorical var -> dummy var
x <- model.matrix(Salary ~ ., Hitters)[-1]
y <- Hitters$Salary
```

### 6.6.1. Ridge Regression

```
# grid for range of values for lambda
grid <- 10^seq(10, -2, length = 100)
# Ridge Regression model
ridge.mod <- glmnet(x, y, alpha = 0, lambda = grid)
# Result is a 20x100 matrix of predictors' coef. depending on lambda
mes <- 'Dimension of result model:'
print(paste(mes, dim(coef(ridge.mod))[1], 'x', dim(coef(ridge.mod))[2]))
```

```
## [1] "Dimension of result model: 20 x 100"
```

```
mes <- 'Available attributes from the model:'
print(cat(mes, '\n', names(ridge.mod), '\n'))
```

```
## Available attributes from the model:
## a0 beta df dim lambda dev.ratio nulldev npasses jerr offset call nob
## NULL
```

Example:

```
# Ex: at index 50, lambda = 11,498
lambda <- ridge.mod$lambda[50]
# l2 norm
l2 <- round(sqrt(sum(coef(ridge.mod)[-1,50]^2)),2)
```

At  $\lambda = 11,497.57$ ,  $l_2$  norm = 6.36, and the coefficients are:

```
coef(ridge.mod)[,50]
```

## (Intercept)	AtBat	Hits	HmRun	Runs
## 407.356050200	0.036957182	0.138180344	0.524629976	0.230701523
## RBI	Walks	Years	CAtBat	CHits
## 0.239841459	0.289618741	1.107702929	0.003131815	0.011653637
## CHmRun	CRuns	CRBI	CWalks	LeagueN
## 0.087545670	0.023379882	0.024138320	0.025015421	0.085028114
## DivisionW	PutOuts	Assists	Errors	NewLeagueN
## -6.215440973	0.016482577	0.002612988	-0.020502690	0.301433531



```
# Ex: at index 50, lambda = 11,498
lambda <- ridge.mod$lambda[60]
# l2 norm
l2 <- round(sqrt(sum(coef(ridge.mod)[-1,60]^2)),2)
```

Compare against different  $\lambda$ :

At  $\lambda = 705.48$ ,  $l_2$  norm = 57.11, and the coefficients are:

```
coef(ridge.mod)[,60]
```

## (Intercept)	AtBat	Hits	HmRun	Runs
## 54.32519950	0.11211115	0.65622409	1.17980910	0.93769713
## RBI	Walks	Years	CAtBat	CHits
## 0.84718546	1.31987948	2.59640425	0.01083413	0.04674557
## CHmRun	CRuns	CRBI	CWalks	LeagueN
## 0.33777318	0.09355528	0.09780402	0.07189612	13.68370191
## DivisionW	PutOuts	Assists	Errors	NewLeagueN
## -54.65877750	0.11852289	0.01606037	-0.70358655	8.61181213

## Predict

```
# predict()
predict(ridge.mod, s = 50, type = 'coefficients')[1:20,]
```

## (Intercept)	AtBat	Hits	HmRun	Runs
## 4.876610e+01	-3.580999e-01	1.969359e+00	-1.278248e+00	1.145892e+00
## RBI	Walks	Years	CAtBat	CHits
## 8.038292e-01	2.716186e+00	-6.218319e+00	5.447837e-03	1.064895e-01
## CHmRun	CRuns	CRBI	CWalks	LeagueN
## 6.244860e-01	2.214985e-01	2.186914e-01	-1.500245e-01	4.592589e+01
## DivisionW	PutOuts	Assists	Errors	NewLeagueN
## -1.182011e+02	2.502322e-01	1.215665e-01	-3.278600e+00	-9.496680e+00

```
# Split into train and test sets
set.seed(1)
train <- sample(1:nrow(x), nrow(x)/2)
test <- (-train)
y.test <- y[test]
```

## Models

```
ridge.mod <- glmnet(x[train,], y[train], alpha = 0, lambda = grid, thresh = 1e-12)
# lambda = s = 4, to specify newdata: newx = x[test,]
ridge.pred <- predict(ridge.mod, s = 4, newx = x[test,])
# MSE
mse.1 <- mean((ridge.pred - y.test)^2)
# MSE for model of only the intercept and no var.: horizontal line y = ybar
mse.2 <- mean((mean(y[train]) - y.test)^2)
# lambda = 1010 aka very large
ridge.pred <- predict(ridge.mod, s = 1e10, newx = x[test,])
```


```
# MSE
mse.3 <- mean((ridge.pred - y.test)^2)
# lambda = 0 <=> Least Squares Regression <=> lm()
ridge.pred <- predict(ridge.mod, s = 0, newx = x[test,])
# MSE
mse.4 <- mean((ridge.pred - y.test)^2)

attr.names <- c('lambda = 4', 'Intercept only', 'lambda = 10^10', 'lambda = 0')
attr.values <- c(mse.1, mse.2, mse.3, mse.4)
data.frame(Models = attr.names, MSE = attr.values)
```

```
##           Models      MSE
## 1      lambda = 4 101036.8
## 2 Intercept only 193253.1
## 3 lambda = 10^10 193253.1
## 4      lambda = 0 114723.6
```

### Cross-Validation

```
par(mar=c(4,4,2,1))
# Cross-Validation: cv.glmnet()
set.seed(1)
cv.out <- cv.glmnet(x[train,], y[train], alpha = 0)
plot(cv.out)
```



```
# Best lambda <=> lambda with min(MSE)
bestlam <- round(cv.out$lambda.min, 2)
# predict() and test set MSE based on best lambda
ridge.pred <- predict(ridge.mod, s = bestlam, newx = x[test,])
mse <- round(mean((ridge.pred - y.test)^2), 0)
```

From Cross-Validation under **Ridge Regression**,  $\lambda$  with the lowest MSE on test set is 211.74, with  $MSE = 9.6016 \times 10^4$ .


```
# Model on the entire dataset
out <- glmnet(x, y, alpha = 0)
predict(out, type = 'coefficients', s = bestlam)[1:20,]
```

##	(Intercept)	AtBat	Hits	HmRun	Runs
##	9.88487135	0.03143885	1.00883177	0.13926743	1.11320858
##	RBI	Walks	Years	CAtBat	CHits
##	0.87318972	1.80410571	0.13071797	0.01113977	0.06489859
##	CHmRun	CRuns	CRBI	CWalks	LeagueN
##	0.45158636	0.12900078	0.13737743	0.02908517	27.18236859
##	DivisionW	PutOuts	Assists	Errors	NewLeagueN
##	-91.63431942	0.19149294	0.04254564	-1.81245301	7.21203276

---


## 6.6.2. The Lasso

```
par(mar=c(4,4,2,1))
# glmnet() with alpha = 1
lasso.mod <- glmnet(x[train,], y[train], alpha = 1, lambda = grid)
plot(lasso.mod)
```



Cross-Validation: `cv.glmnet()`

```
# Cross-Validation: cv.glmnet()
par(mar=c(4,4,2,1))
set.seed(1)
cv.out <- cv.glmnet(x[train,], y[train], alpha = 1)
plot(cv.out)
```



```
# Best lambda <=> lambda with min(MSE)
bestlam <- cv.out$lambda.min
# predict() and test set MSE based on best lambda
lasso.pred <- predict(lasso.mod, s = bestlam, newx = x[test,])
mse <- mean((lasso.pred - y.test)^2)
```

From Cross-Validation under **Lasso**,  $\lambda$  with the lowest MSE on test set is 16.7801585, with  $MSE = 1.0074345 \times 10^5$ .

```
# Model on the entire dataset
out <- glmnet(x, y, alpha = 1, lambda = grid)
predict(out, type = 'coefficients', s = bestlam)[1:20,]
```

## (Intercept)	AtBat	Hits	HmRun	Runs
## 18.5394844	0.0000000	1.8735390	0.0000000	0.0000000
## RBI	Walks	Years	CAtBat	CHits
## 0.0000000	2.2178444	0.0000000	0.0000000	0.0000000
## CHmRun	CRuns	CRBI	CWalks	LeagueN
## 0.0000000	0.2071252	0.4130132	0.0000000	3.2666677
## DivisionW	PutOuts	Assists	Errors	NewLeagueN
## -103.4845458	0.2204284	0.0000000	0.0000000	0.0000000

#### Comment:

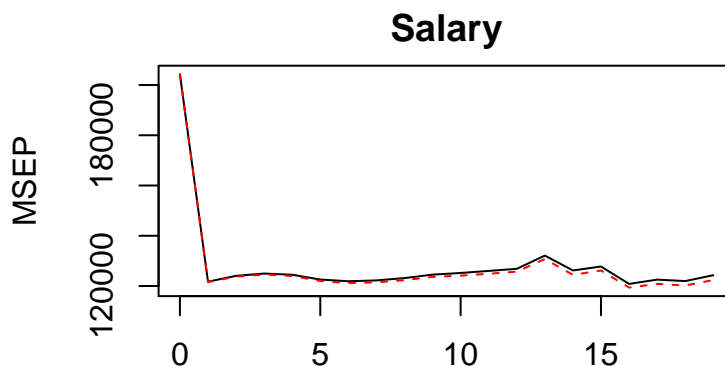
- Lasso and Ridge regression models' MSE are similar
- Lasso model: 12 vars have coef = 0 => Use fewer vars => Easier to interpret

## 6.7. Lab 3: PCR and PLS Regression

### 6.7.1. Principal Components Regression

`pcr()` from `pls` package, > `scale = TRUE` to normalize data, > `validation = 'CV'` for 10-fold Cross Validation by default.


```
# `pcr()` `` fn from `pls` `` package
library(pls)
par(mar=c(2,4,2,1))
set.seed(2)
pcr.fit <- pcr(Salary ~ ., data = Hitters, scale = TRUE, validation = 'CV')
# summary(pcr.fit)
# Plot: `validationplot()`, to specify MSE: `val.type = 'MSEP'`
validationplot(pcr.fit, val.type = 'MSEP')
```



Comment:

- printed are RMSE, to get  $MSE = RMSE^2$
- CV MSE is smallest when  $M = 16$ , not much different from  $M = 19 \iff$  no reduction
- `summary()` shows Percentage of Variance Explained

```
par(mar=c(2,4,2,1))
# Model from train and test sets
set.seed(1)
pcr.fit <- pcr(Salary ~ ., data = Hitters, subset = train, scale = TRUE, validation = 'CV')
validationplot(pcr.fit, val.type = 'MSEP')
```



```
# Prediction, based on M = 7
pcr.pred <- predict(pcr.fit, x[test,], ncomp = 7)
mean((pcr.pred - y.test)^2)
```

```
## [1] 96556.22
```

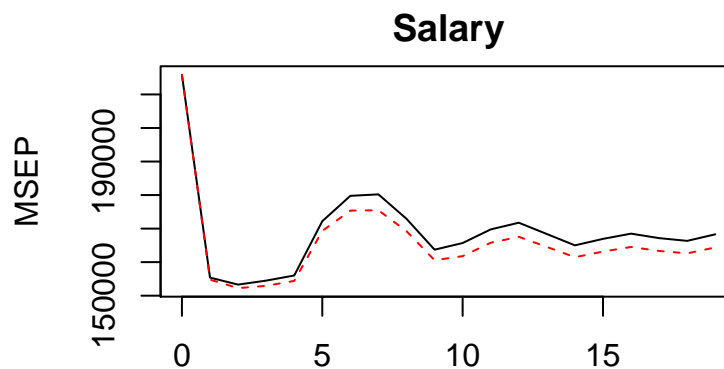
**Comment:** MSE is competitive vs. Ridge Regression and the Lasso, However, model from PCR is more difficult to interpret.

```
# Fit model on full data set, using M = 7
pcr.fit <- pcr(y ~ x, scale = TRUE, ncomp = 7)
# summary(pcr.fit)
```

## 6.7.2. Partial Least Squares

`pls()` from `pls` package

```
par(mar=c(2,4,2,1))
set.seed(1)
pls.fit <- pls(Salary ~ ., data = Hitters, subset = train, scale = TRUE, validation = 'CV')
#summary(pls.fit)
# Plot
validationplot(pls.fit, val.type = 'MSEP')
```



```
pls.pred <- predict(pls.fit, x[test,], ncomp = 2)
mean((pls.pred - y.test)^2)
```

```
## [1] 101417.5
```

**Comment:** MSE is comparable but slightly higher than Ridge Regression, the Lasso, and PCR

```
# Fit model on full data set, using M = 2
pls.fit <- pls(Salary ~ ., data = Hitters, scale = TRUE, ncomp = 2)
summary(pls.fit)
```

```
## Data:      X dimension: 263 19
## Y dimension: 263 1
## Fit method: kernelpls
## Number of components considered: 2
```

```
## TRAINING: % variance explained
##          1 comps  2 comps
## X          38.08   51.03
## Salary     43.05   46.40
```

**Comment:** PLSR model with 2 components explains 46.40% variance in Salary while PCR needs 7 components to explain 46.69 %%

**Reason:** PCR only attempts to maximize variance explained in the predictors while PLSR searches for Directions that explain the variance in both predictors and response.